

In [1]:

```
#student_name:YonghengHou
#student_number:5556661
#login:yh790
from pyspark.sql import SparkSession
import numpy as np
```

In [2]:

```
spark = SparkSession.builder.appName("CSCI316-project") \
    .config("spark-master", "local") \
    .getOrCreate()
spark
```

Out[2]:

SparkSession - in-memory
SparkContext

[Spark UI \(http://windows10.microdone.cn:4041\)](http://windows10.microdone.cn:4041)

Version

v2.4.4

Master

local[*]

AppName

CSCI316-project

In [3]:

```
from pyspark.sql.types import *
myManualSchema = StructType([
    StructField("Session_ID", IntegerType(), True),
    StructField("Timestamp", TimestampType(), True),
    StructField("Item_ID", IntegerType(), True),
    StructField("Category", StringType(), True)])

df_fd = spark \
    .read \
    .format("csv") \
    .schema(myManualSchema).load("yoochoose-clicks.dat")
```

In [4]:

```
df_fd.printSchema()
```

```
root
|-- Session_ID: integer (nullable = true)
|-- Timestamp: timestamp (nullable = true)
|-- Item_ID: integer (nullable = true)
|-- Category: string (nullable = true)
```

In [5]:

```
#transfer the 8-10 digits number to "B"
from pyspark.sql.functions import regexp_replace,col
cateFilter="[0-9]{7,10}"
df_FD=df_fd.withColumn("Category", regexp_replace(col("Category"), cateFilter, "B"))
df_FD.show()
```

Session_ID	Timestamp	Item_ID	Category
1	2014-04-07 20:51:...	214536502	0
1	2014-04-07 20:54:...	214536500	0
1	2014-04-07 20:54:...	214536506	0
1	2014-04-07 20:57:...	214577561	0
2	2014-04-07 23:56:...	214662742	0
2	2014-04-07 23:57:...	214662742	0
2	2014-04-07 23:58:...	214825110	0
2	2014-04-07 23:59:...	214757390	0
2	2014-04-08 00:00:...	214757407	0
2	2014-04-08 00:02:...	214551617	0
3	2014-04-03 00:17:...	214716935	0
3	2014-04-03 00:26:...	214774687	0
3	2014-04-03 00:30:...	214832672	0
4	2014-04-07 22:09:...	214836765	0
4	2014-04-07 22:26:...	214706482	0
6	2014-04-07 02:58:...	214701242	0
6	2014-04-07 03:02:...	214826623	0
7	2014-04-02 17:38:...	214826835	0
7	2014-04-02 17:39:...	214826715	0
8	2014-04-06 18:49:...	214838855	0

only showing top 20 rows

In [6]:

```
#create second same dataframe and change column name in order to join this two tables in the next steps
second_df=df_FD.selectExpr("Session_ID as newSession_ID","Timestamp as newTimestamp","Item_ID as newItem_ID","Category as newCategory")

#for join, I index two dataframes, and one dataframe index from 0, another one is index from 1,
#it can has dislocation by subtracting to calculate interval time
from pyspark.sql.functions import monotonically_increasing_id
df1 = df_FD.withColumn("index", monotonically_increasing_id()+1)
df2 =second_df.withColumn("index", (monotonically_increasing_id()))
df1.show(10)
df2.show(10)
```

Session_ID	Timestamp	Item_ID	Category	index
1	2014-04-07 20:51:...	214536502	0	1
1	2014-04-07 20:54:...	214536500	0	2
1	2014-04-07 20:54:...	214536506	0	3
1	2014-04-07 20:57:...	214577561	0	4
2	2014-04-07 23:56:...	214662742	0	5
2	2014-04-07 23:57:...	214662742	0	6
2	2014-04-07 23:58:...	214825110	0	7
2	2014-04-07 23:59:...	214757390	0	8
2	2014-04-08 00:00:...	214757407	0	9
2	2014-04-08 00:02:...	214551617	0	10

only showing top 10 rows

newSession_ID	newTimestamp	newItem_ID	newCategory	index
1	2014-04-07 20:51:...	214536502	0	0
1	2014-04-07 20:54:...	214536500	0	1
1	2014-04-07 20:54:...	214536506	0	2
1	2014-04-07 20:57:...	214577561	0	3
2	2014-04-07 23:56:...	214662742	0	4
2	2014-04-07 23:57:...	214662742	0	5
2	2014-04-07 23:58:...	214825110	0	6
2	2014-04-07 23:59:...	214757390	0	7
2	2014-04-08 00:00:...	214757407	0	8
2	2014-04-08 00:02:...	214551617	0	9

only showing top 10 rows

In [7]:

```
#lifter join two tables, df1 and df2
df3 = df1.join(df2, "index", "left_outer")
df3.show(10)
```

index	Session_ID	Timestamp	Item_ID	Category	newSession_ID	newTimestamp	newItem_ID	newCategory
26	11	2014-04-03 21:45:...	214821275	0	11	2014-04-03 21:45:...	214821371	0
29	11	2014-04-03 21:46:...	214821371	0	11	2014-04-03 21:53:...	214717089	0
474	154	2014-04-03 20:04:...	214560187	0	154	2014-04-03 20:05:...	214716984	0
964	337	2014-04-04 05:52:...	214820842	0	337	2014-04-04 05:56:...	214826897	0
1677	564	2014-04-02 21:49:...	214629060	0	564	2014-04-02 22:09:...	214840899	0
1697	564	2014-04-02 23:26:...	214596647	0	564	2014-04-02 23:27:...	214837558	0
1806	531	2014-04-02 01:58:...	214748336	0	531	2014-04-02 01:59:...	214717247	0
1950	638	2014-04-02 23:37:...	214579730	0	637	2014-04-02 05:54:...	214537867	0
2040	603	2014-04-07 18:20:...	214684513	0	602	2014-04-02 23:21:...	214819562	0
2214	661	2014-04-02 04:40:...	214832559	0	661	2014-04-02 04:41:...	214819550	0

only showing top 10 rows

In [8]:

```
#this function is to calculate interval time and put the result into new column
import pyspark.sql.functions as F
df4=df3.withColumn(
    "interval_time",
    (F.col("newTimestamp").cast("long") - F.col("Timestamp").cast("long")))
df4.show(10)
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
|index|Session_ID|Timestamp|Item_ID|Category|newSession_ID|newTimestamp|newItem_ID|newCategory|interval_time|
+-----+-----+-----+-----+-----+-----+-----+
| 26|11|2014-04-03 21:45:...|214821275|0|11|2014-04-03 21:45:...|214821371|0|28|
| 29|11|2014-04-03 21:46:...|214821371|0|11|2014-04-03 21:53:...|214717089|0|385|
| 474|154|2014-04-03 20:04:...|214560187|0|154|2014-04-03 20:05:...|214716984|0|42|
| 964|337|2014-04-04 05:52:...|214820842|0|337|2014-04-04 05:56:...|214826897|0|222|
| 1677|564|2014-04-02 21:49:...|214629060|0|564|2014-04-02 22:09:...|214840899|0|1226|
| 1697|564|2014-04-02 23:26:...|214596647|0|564|2014-04-02 23:27:...|214837558|0|49|
| 1806|531|2014-04-02 01:58:...|214748336|0|531|2014-04-02 01:59:...|214717247|0|23|
| 1950|638|2014-04-02 23:37:...|214579730|0|637|2014-04-02 05:54:...|214537867|0|-63803|
| 2040|603|2014-04-07 18:20:...|214684513|0|602|2014-04-02 23:21:...|214819562|0|-417493|
| 2214|661|2014-04-02 04:40:...|214832559|0|661|2014-04-02 04:41:...|214819550|0|11|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+
only showing top 10 rows
```

In [9]:

```
#because all session_id is continuous, so it will come out the condition that last record of one session_id
# subtract the first record of another one session_id, so it is wrong condition, so following step is to filter this
"""
example:
last column negative number

/ 1950/      638/2014-04-02 23:37:... /214579730/      0/      637/2014-04-02 05:54:... / 21
4537867/      0/      -63803/
/ 2040/      603/2014-04-07 18:20:... /214684513/      0/      602/2014-04-02 23:21:... / 21
4819562/      0/      -417493/
/ 2214/      661/2014-04-02 04:40:... /214832559/      0/      661/2014-04-02 04:41:... / 21
4819550/      0/      11/

"""
df5=df4.filter((col("Session_ID") == col("newSession_ID")))
df5.show(10)
```

index	Session_ID	Timestamp	Item_ID	Category	newSession_ID	newTimestamp	newItem_ID	newCategory	interval_time
35	11	2014-04-03 21:57:...	214826837	0	11	2014-04-03 21:57:...	214819762	0	20
201	62	2014-04-07 01:44:...	214826619	0	62	2014-04-07 01:45:...	214746427	0	53
288	86	2014-04-02 05:21:...	214648340	0	86	2014-04-02 05:21:...	214648438	0	31
643	219	2014-04-01 18:07:...	214725500	0	219	2014-04-01 18:10:...	214839660	0	161
1219	389	2014-04-07 03:20:...	214691396	0	389	2014-04-07 03:21:...	214691321	0	42
1184	397	2014-04-05 18:19:...	214553540	0	397	2014-04-05 18:19:...	214572538	0	26
1310	479	2014-04-03 20:09:...	214820814	0	479	2014-04-03 20:11:...	214746339	0	144
1521	496	2014-04-08 03:39:...	214638977	0	496	2014-04-08 03:39:...	214638977	0	9
1647	554	2014-04-08 11:36:...	214829312	0	554	2014-04-08 11:36:...	214774685	0	33
2154	651	2014-04-02 04:17:...	214718117	0	651	2014-04-02 04:17:...	214690845	0	37

only showing top 10 rows

In [10]:

```
#1. this step is for the condition that in one session_id, it has more than or equal two same item_id operations,
#because this operations belong to one item, so group it, and sum the interval_time of same item_id records.
#2. why I max("Category") here, because I just want to keep Category column after groupby, But I didnt find good
#solution to generate it, so I have to do it like this.
"""
example: interval time of index 4 and index 5 sum together
            item_id           cate  index
2/2014-04-07 23:56:... / 214662742/      0/    4/
/      2/2014-04-07 23:57:... / 214662742/      0/    5/
/      2/2014-04-07 23:58:... / 214825110/      0/    6/
/      2/2014-04-07 23:59:... / 214757390/      0/    7/
/      2/2014-04-08 00:00:... / 214757407/      0/    8/
/      2/2014-04-08 00:02:... / 214551617/      0/    9/
"""

df6=df5.groupby([col("Session_ID"),col("Item_ID")])\
.agg(F.sum(col("interval_time")).alias("complete_interval_time"),F.max(col("Category")).alias("Category"))
df6.show(10)
```

Session_ID	Item_ID	complete_interval_time	Category
87	214554637	13	0
119	214716954	794	0
397	214843492	127	0
491	214832559	42	0
492	214718220	93	0
516	214676364	59	0
577	214708372	48	0
626	214827005	54	0
651	214838833	252	0
726	214589632	23	0

only showing top 10 rows

In [11]:

```
#1. so far, we have the complete_interval_time for each user on each same item
# 2, group by category, then sum the complete_interval_time, then divide by total number of each user each item
df7=df6.groupby([col("Category")]).agg(F.count(col('Item_ID')),F.sum(col('complete_interval_time')))
df7.show(10)
```

Category	count(Item_ID)	sum(complete_interval_time)
7	220258	47393374
11	44447	6591571
3	534948	70182484
8	27314	5099778
0	9850312	1725863516
5	257109	62751734
B	46800	9880439
6	223263	52894611
S	6534387	1111613782
9	61836	12428756

only showing top 10 rows

In [12]:

```
#lastly, calculate the average time
final_result=df7.withColumn("Average_time", (col("sum(complete_interval_time)"/col("count(Item_ID)"))))
final_result.show(15)
```

Category	count(Item_ID)	sum(complete_interval_time)	Average_time
7	220258	47393374	215.17208909551525
11	44447	6591571	148.30182014534165
3	534948	70182484	131.1949647442368
8	27314	5099778	186.70930658270484
0	9850312	1725863516	175.20902038432894
5	257109	62751734	244.0666565542241
B	46800	9880439	211.12049145299144
6	223263	52894611	236.91615269883502
S	6534387	1111613782	170.117530841072
9	61836	12428756	200.99547189339543
1	1013746	201798049	199.0617462362367
10	41407	8073943	194.9898084864878
4	281139	55556414	197.61190727718318
12	9672	2787961	288.25072373862696
2	742758	157088599	211.49364799840595

In []: